

# HW7

Melchor Ronquillo

2022-12-08

8. This problem involves the OJ data set which is part of the ISLR2 package.

- (a) Create a training set containing a random sample of 800 observations, and a test set containing the remaining observations.

```
install.packages('ISLR2', repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/b4/vbzhgztj3tj299xgxnlkp28c0000gn/T//RtmpFbKzMC/downloaded_packages
library(ISLR2)

Juice <- OJ
#Juice

set.seed(777)
randomJuice= Juice[sample(1:nrow(Juice)), ]
#randomJuice

nrow(randomJuice)

## [1] 1070

train_OJ <- randomJuice[1:800,]
test_OJ <- randomJuice[801:1070,]

#train_OJ
#test_OJ
```

- (b) Fit a support vector classifier to the training data using cost = 0.01, with Purchase as the response and the other variables as predictors. Use the summary() function to produce summary statistics, and describe the results obtained.

```
install.packages('e1071', repos = "http://cran.us.r-project.org")

##
## The downloaded binary packages are in
## /var/folders/b4/vbzhgztj3tj299xgxnlkp28c0000gn/T//RtmpFbKzMC/downloaded_packages
library(e1071)

set.seed(778)
JuiceSVC <- svm(Purchase ~ ., data = train_OJ, kernel = "linear", cost = 0.01, scale = TRUE)
#plot(JuiceSVC, train_OJ)
summary(JuiceSVC)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = train_OJ, kernel = "linear", cost = 0.01,
##      scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##      cost:  0.01
##
## Number of Support Vectors:  432
##
## ( 216 216 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
```

(c) What are the training and test error rates?

```
train_OJ_pred = predict(JuiceSVC, train_OJ)
table(train_OJ$Purchase, train_OJ_pred)
```

```
##      train_OJ_pred
##      CH  MM
## CH 433  57
## MM  75 235
```

```
train_error = (57+75) / (453+57+75+235)
train_error
```

```
## [1] 0.1609756
```

```
      Train error rate:
      16% error
```

```
test_OJ_pred = predict(JuiceSVC, test_OJ)
table(test_OJ$Purchase, test_OJ_pred)
```

```
##      test_OJ_pred
##      CH  MM
## CH 137  26
## MM  24  83
```

```
test_error = (26+24) / (137+26+24+83)
test_error
```

```
## [1] 0.1851852
```

```
      Test error rate:
      18.5%
```

(d) Use the tune() function to select an optimal cost. Consider values in the range 0.01 to 10.

```

set.seed(779)
tuneJuice <- tune(svm, Purchase ~ ., data = Juice, kernel = "linear",
                 ranges = list(cost = c(0.01, 0.05, 0.1, 0.5, 1, 5, 10))
                 )
summary(tuneJuice)

```

```

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##     1
##
## - best performance: 0.1654206
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.1691589 0.03098083
## 2  0.05 0.1700935 0.03261371
## 3  0.10 0.1710280 0.03470411
## 4  0.50 0.1682243 0.03296886
## 5  1.00 0.1654206 0.03442332
## 6  5.00 0.1672897 0.03002637
## 7 10.00 0.1700935 0.03170844

```

(e) Compute the training and test error rates using this new value for cost.

```

bestmod <- tuneJuice$best.model
summary(bestmod)

```

```

##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ ., data = Juice, ranges = list(cost = c(0.01,
##   0.05, 0.1, 0.5, 1, 5, 10)), kernel = "linear")
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: linear
##     cost:  1
##
## Number of Support Vectors:  442
##
## ( 221 221 )
##
##
## Number of Classes:  2
##
## Levels:
##   CH MM

```

```

bestpred_OJ_SVC <- predict(bestmod, test_OJ)
table(test_OJ$Purchase, bestpred_OJ_SVC)

```

```
##      bestpred_OJ_SVC
##      CH  MM
##  CH 140  23
##  MM  23  84

test_error = (23+23) / (140+23+23+84)
test_error
```

```
## [1] 0.1703704

      Test Error: 17%
```

(f) Repeat parts (b) through (e) using a support vector machine with a radial kernel. Use the default value for gamma.

```
#Create SVM
set.seed(780)
JuiceSVM <- svm(Purchase ~ ., data = train_OJ, kernel = "radial",
                gamma = 1, cost = 0.01, scale = TRUE)
#plot(JuiceSVM, train_OJ)
summary(JuiceSVM)

##
## Call:
## svm(formula = Purchase ~ ., data = train_OJ, kernel = "radial", gamma = 1,
##      cost = 0.01, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: radial
##      cost:  0.01
##
## Number of Support Vectors:  653
##
##  ( 343 310 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM

#train and test errors
train_OJ_predSVM = predict(JuiceSVM, train_OJ)
table(train_OJ$Purchase, train_OJ_predSVM)

##      train_OJ_predSVM
##      CH  MM
##  CH 490   0
##  MM 310   0

test_OJ_predSVM = predict(JuiceSVM, test_OJ)
table(test_OJ$Purchase, test_OJ_predSVM)

##      test_OJ_predSVM
##      CH  MM
```

```

## CH 163 0
## MM 107 0

#Tune
set.seed(781)
tuneJuiceSVM <- tune(svm, Purchase ~ ., data = Juice, kernel = "radial",
                    ranges = list(cost = c(0.01, 0.05, 0.1, 0.5, 1, 5, 10))
                    )

summary(tuneJuiceSVM)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
## cost
## 0.5
##
## - best performance: 0.1728972
##
## - Detailed performance results:
## cost error dispersion
## 1 0.01 0.3897196 0.04746052
## 2 0.05 0.1971963 0.02222562
## 3 0.10 0.1869159 0.02530853
## 4 0.50 0.1728972 0.03475999
## 5 1.00 0.1728972 0.02963598
## 6 5.00 0.1738318 0.03152426
## 7 10.00 0.1766355 0.03002637

bestmodSVM <- tuneJuiceSVM$best.model
summary(bestmodSVM)

##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ ., data = Juice, ranges = list(cost = c(0.01,
## 0.05, 0.1, 0.5, 1, 5, 10)), kernel = "radial")
##
##
## Parameters:
## SVM-Type: C-classification
## SVM-Kernel: radial
## cost: 0.5
##
## Number of Support Vectors: 523
##
## ( 265 258 )
##
##
## Number of Classes: 2
##
## Levels:
## CH MM

```

```
bestpred_OJ_SVM <- predict(bestmodSVM, test_OJ)
table(test_OJ$Purchase, bestpred_OJ_SVM)
```

```
##      bestpred_OJ_SVM
##      CH  MM
## CH 149  14
## MM  26  81
```

```
test_error = (14+26) / (149+14+26+81)
test_error
```

```
## [1] 0.1481481
```

```
Test Error: 14.8%
```

(g) Repeat parts (b) through (e) using a support vector machine with a polynomial kernel. Set degree = 2.

```
set.seed(782)
JuiceSVCPoly <- svm(Purchase ~ ., data = train_OJ, kernel = "polynomial",
                    cost = 0.01,
                    degree = 2,
                    scale = TRUE)
summary(JuiceSVCPoly)
```

```
##
## Call:
## svm(formula = Purchase ~ ., data = train_OJ, kernel = "polynomial",
##      cost = 0.01, degree = 2, scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##      cost:   0.01
##   degree:    2
##   coef.0:    0
##
## Number of Support Vectors:  627
##
## ( 317 310 )
##
##
## Number of Classes:  2
##
## Levels:
## CH MM
```

```
#train and test errors
train_OJ_predPoly = predict(JuiceSVCPoly, train_OJ)
table(train_OJ$Purchase, train_OJ_predPoly)
```

```
##      train_OJ_predPoly
##      CH  MM
```

```

##    CH 489    1
##    MM 306    4

test_OJ_predPoly = predict(JuiceSVCPoly, test_OJ)
table(test_OJ$Purchase, test_OJ_predPoly)

##      test_OJ_predPoly
##      CH  MM
##    CH 163    0
##    MM 106    1

#tune
set.seed(779)
tuneJuicePoly <- tune(svm, Purchase ~ ., data = Juice, kernel = "polynomial",
                      degree = 2, scale = TRUE,
                      ranges = list(cost = c(0.01, 0.05, 0.1, 0.5, 1, 5, 10))
                      )
summary(tuneJuicePoly)

##
## Parameter tuning of 'svm':
##
## - sampling method: 10-fold cross validation
##
## - best parameters:
##   cost
##   10
##
## - best performance: 0.1757009
##
## - Detailed performance results:
##   cost      error dispersion
## 1  0.01 0.3691589 0.04322254
## 2  0.05 0.3299065 0.04340180
## 3  0.10 0.2990654 0.04427623
## 4  0.50 0.2056075 0.02643390
## 5  1.00 0.1943925 0.02636037
## 6  5.00 0.1785047 0.02970140
## 7 10.00 0.1757009 0.03935604

#best model
bestmodPoly <- tuneJuicePoly$best.model
summary(bestmodPoly)

##
## Call:
## best.tune(METHOD = svm, train.x = Purchase ~ ., data = Juice, ranges = list(cost = c(0.01,
##   0.05, 0.1, 0.5, 1, 5, 10)), kernel = "polynomial", degree = 2,
##   scale = TRUE)
##
##
## Parameters:
##   SVM-Type:  C-classification
##   SVM-Kernel: polynomial
##   cost:      10
##   degree:    2

```

```
##      coef.0:  0
##
## Number of Support Vectors:  450
##
## ( 230 220 )
##
##
## Number of Classes:  2
##
## Levels:
##  CH MM
bestpred_OJ_Poly <- predict(bestmodPoly, test_OJ)
table(test_OJ$Purchase, bestpred_OJ_Poly)
```

```
##      bestpred_OJ_Poly
##      CH  MM
##  CH 148  15
##  MM  25  82
test_error = (15+25) / (148+15+25+82)
test_error
```

```
## [1] 0.1481481
```

Test Error: 14.8%

(h) Overall, which approach seems to give the best results on this data?

Overall, the approach with the best results was both radial and polynomial, as they both had the lowest test error rate of 14.8%